

# AI Music Generators Comparison

by Anna Zhdanova

	Suno	Udio	Mubert
Website	<a href="https://suno.com">suno.com</a>	<a href="https://udio.com">udio.com</a>	<a href="https://mubert.com">mubert.com</a>
Subjective quality of the result	<p>Suno can transform simple prompts into full songs that sound professionally produced, with natural-sounding instruments and structure.</p> <p>The tracks often feel polished in structure and harmony.</p> <p>Stands out for its highly realistic and emotionally expressive vocals, especially in the paid version (v4.5)</p> <p>My personal favorite</p>	<p>Udio can produce high-quality audio with strong rhythm and genre fidelity.</p> <p>While the output sounds good on the surface, it lacks emotional depth, originality, and memorable melodies.</p> <p>There are occasional issues with harmonic complexity. Besides, the vocals can sound artificial or overly autotuned.</p> <p>My personal #2</p>	<p>Mubert is designed more for creating high-quality instrumental or ambient background music. It excels at generating customizable soundscapes for content creators, such as for videos, apps, or podcasts.</p> <p>While the sound is generally clean and well-produced, it may lack the dynamic range, emotional feel, or complexity needed for full-fledged songs with vocals.</p>
Unique features compared to other platforms	<ul style="list-style-type: none"><li>• Strong vocal support</li><li>• Lyric-aware generation,</li><li>• Intuitive style inference</li><li>• Full song structure with verses and choruses.</li></ul>	<ul style="list-style-type: none"><li>• High musicality</li><li>• Adapts well to mood and text</li><li>• Excels at emotional delivery by instrumental music and multi-genre stylization.</li></ul>	<ul style="list-style-type: none"><li>• Continuous genre-based music generation</li><li>• API integration</li><li>• Audio streams</li><li>• Music for NFTs or brand use</li></ul>

Can accept band name as a reference?  Recognizes obscure electronic groups?	Yes, recognizes both famous and niche groups. For example, understands 'dark techno like Ancient Methods'. Results are not always exact but stylistically relevant.	Yes, handles artist names and rare genres well (e.g., 'ethereal industrial pop'). Output is often stylistically close.	Mostly no. Works better with genre/atmosphere tags than artist names. 'Cyberpunk techno' yields better results than 'Phase Fatale'.
Does it support audio track as input?	Yes – Pro and Premier users can upload or record audio (6–60 sec, up to 8 min) for continuation or cover generation.	Yes – allows upload of audio up to 15 minutes to generate forward/backward extensions.	Yes. In Mubert Render Pro or Creator Studio, you can upload samples and use them as a generation source.
Maximum generated track length	Pro and Premier users can upload or record audio (6–60 sec, up to 8 min) for continuation or cover generation.	Each segment is ~30 sec, but can be extended to ~15 minutes using multiple parts.	Streaming: infinite. Downloadable renders: up to 25 minutes.
Can you edit specific parts of a generated or uploaded track?	Enables you to highlight any verse, chorus, instrumental riff, etc., and regenerate just that slice — without touching the rest of the composition	Lets you select a region (up to ~28 seconds) within a track and regenerate just that part — ideal for correcting lyrics, changing instrumentation, or smoothing transitions	Partial: allows intro/bridge/outro control. No fine post-editing, but you can manipulate blocks during render.
Subjective ease of use (how easy is it to sign up and create your first track)	Extremely easy: email or Google login, immediate access. First track can be created in under 30 seconds.	Fast signup, instant access after login. User interface is intuitive; track creation takes less than a minute.	Less intuitive interface. First generation takes 2–3 minutes due to more configuration steps.
Date of release	2023	2024	2017

Founders	Michael Shulman, Georg Kucsko, Martin Camacho, Keenan Freyberg. Team has background at Meta, Tiktok, and MIT	David Ding, Conor Durkan, Charlie Nash, Yaroslav Ganin, Andrew Sanchez – all ex-DeepMind researchers	Russian startup founded by Alexander Kochetkov
Technology	<p>Technical explanations are available via <a href="https://sunnoai.com/?utm">https://sunnoai.com/?utm</a></p> <p>Uses a diffusion-based model trained on a large corpus of music, including vocals.</p> <p>Generates music from text descriptions, including style, tempo, and lyrics.</p>	<p>General tech description available via <a href="https://www.udio.com/guide?utm">https://www.udio.com/guide?utm</a></p> <p>Model trained on musical and vocal tracks.</p> <p>Generates based on descriptions, supporting style, genre, emotion, and lyrics.</p>	<p>First-party description is very vague: “Millions of samples from hundreds of artists are used by Mubert to instantly generate royalty-free AI music”</p> <p>Some speculation on how it worked in 2019 <a href="https://medium.com/%40alexbainter/how-mubert-probably-works-e44de23c45bd">https://medium.com/%40alexbainter/how-mubert-probably-works-e44de23c45bd</a> (paywalled)</p> <p>Probabilistic combination of pre-existing samples and streams using Machine Learning</p>
Monetization	Paid plans offer more generations per day (up to 100), faster rendering, private mode, and export to .stem and .wav formats. Pricing: \$8 and \$24/month.	Subscription increases daily track limits, generation length, and access to faster servers. Plans: \$8, \$24, \$60/month.	Paid plans unlock commercial licenses, WAV downloads, and custom rendering. Plans: \$11, \$32, \$149

